



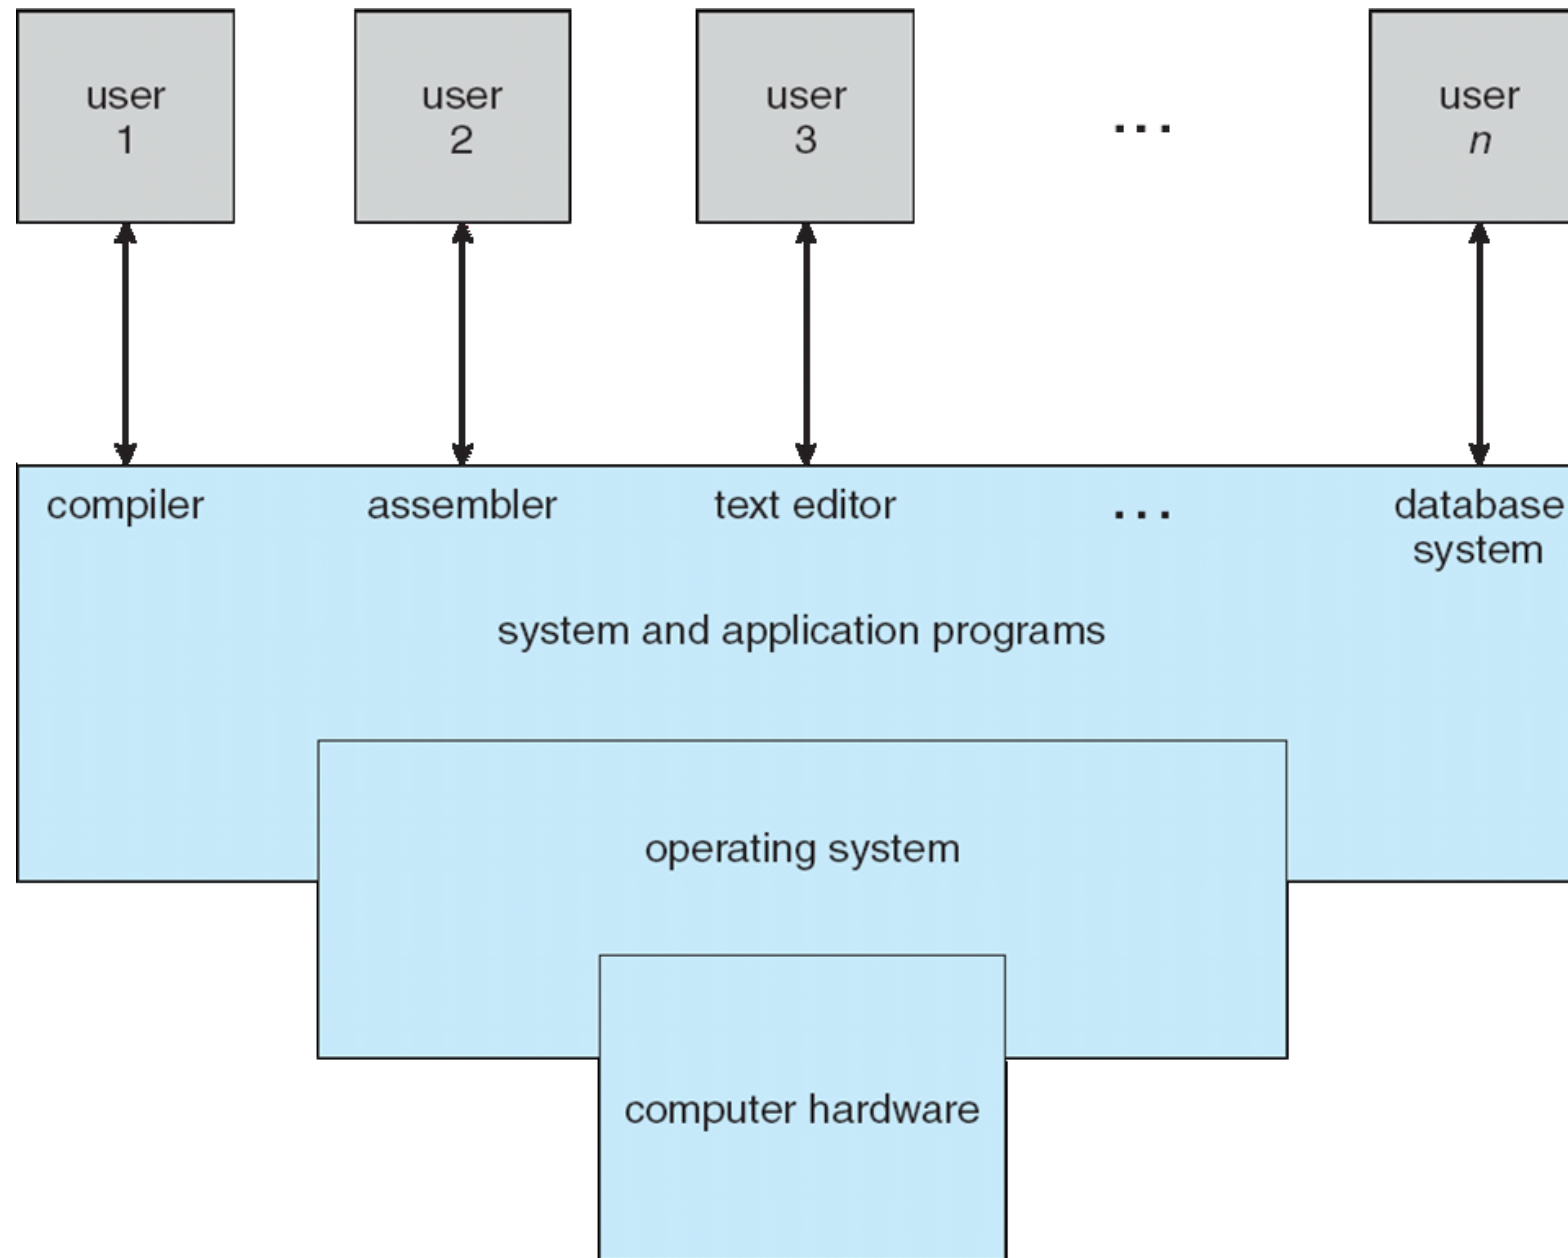
Review 01

Yajin Zhou (<http://yajin.org>)

Zhejiang University

01: Introduction

Four Components of a Computer System





Interrupts and Traps

- Interrupt transfers control to the interrupt service routine
 - **interrupt vector:** a table containing addresses of all the service routines
 - incoming interrupts are disabled while serving another interrupt to prevent a lost interrupt
 - **interrupt handler** must save the (interrupted) execution states
- A **trap** is a **software-generated interrupt**, caused either by an error or a user request
 - an **interrupt** is asynchronous; a **trap** is synchronous
 - e.g., **system call**, divided-by-zero exception, general protection exception...
 - Int 0x80 - this is **not** a privileged instruction
- Operating systems are usually **interrupt-driven**



Interrupt Handling

- Operating system preserves the execution state of the CPU
 - save registers and the program counter (PC)
- OS determines which device caused the interrupt
 - **polling**
 - **vectored** interrupt system
- OS handles the interrupt by calling the device's driver
- OS restores the CPU execution to the saved state



I/O: from System Call to Devices, and Back

- A program uses a **system call** to access system resources
 - e.g., files, network
- Operating system converts it to device access and issues I/O requests
 - I/O requests are sent to the device driver, then to the controller
 - e.g., read disk blocks, send/receive packets...
- OS puts the program to wait (**synchronous I/O**) or returns to it without waiting (**asynchronous I/O**)
 - OS may switch to another program when the requester is waiting
- I/O completes and the controller **interrupts** the OS
- OS processes the I/O, and then wakes up the program (synchronous I/O) or send its a signal (asynchronous I/O)



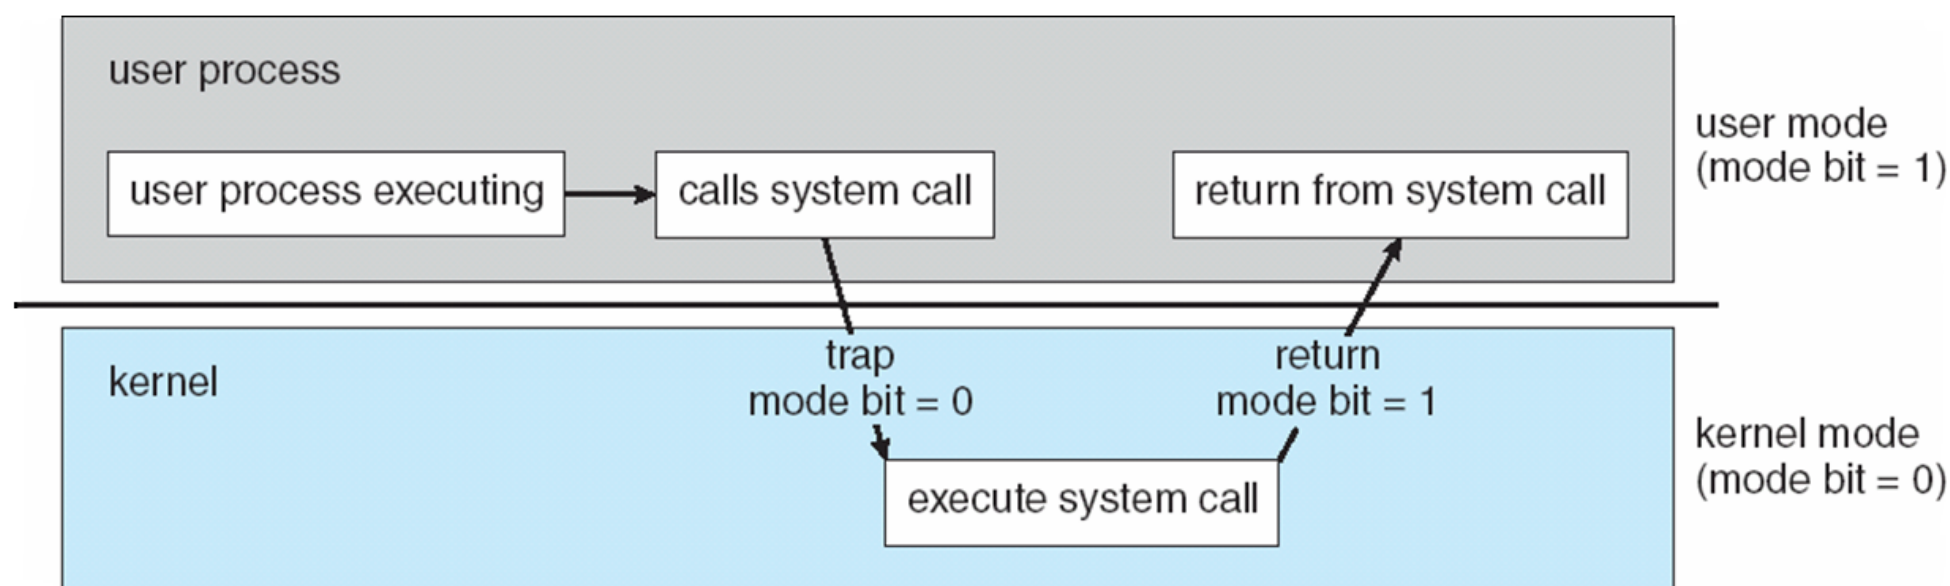
Dual-mode operation

- Operating system is usually interrupt-driven (why?)
 - Efficiency, regain control (timer interrupt)
- **Dual-mode operation** allows OS to protect itself and other system components
 - **user mode** and **kernel mode (or other names)**
 - a **mode** bit distinguishes when CPU is running user code or kernel code
 - some instructions designated as **privileged, only executable in kernel and cannot executed in user mode (and certain memory cannot be accessed in user mode!)**
 - **system call** changes mode to kernel, return from call resets it to user



Transition between Modes

- **System calls, exceptions, interrupts** cause transitions between kernel/user modes



In user mode, certain instructions cannot be executed

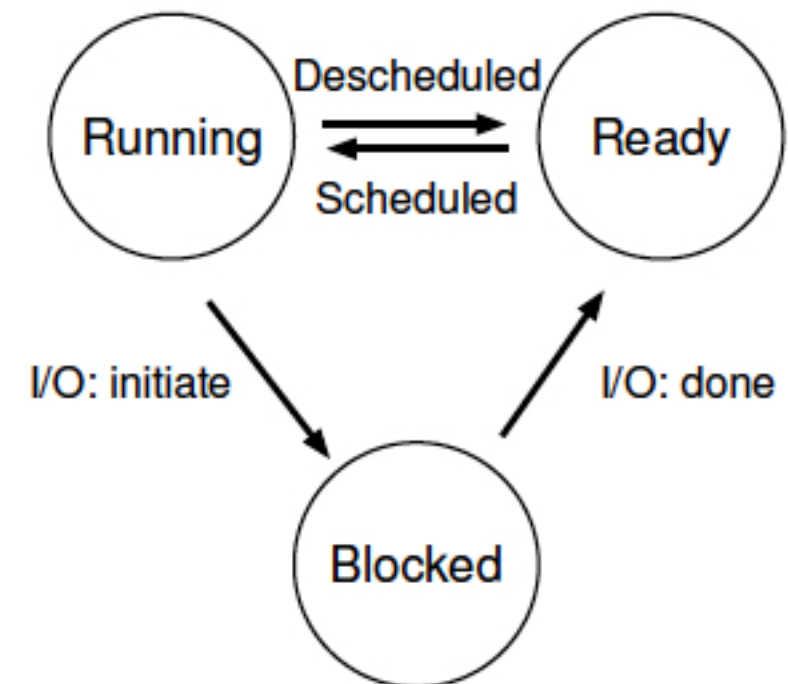


Resource Management: Process Management

- A process is **a program in execution**
 - program is a ***passive*** entity, process is an ***active*** entity
 - a system has many processes running concurrently
- Process needs resources to accomplish its task
 - OS reclaims all reusable resources upon process termination
 - e.g., CPU, memory, I/O, files, initialization data


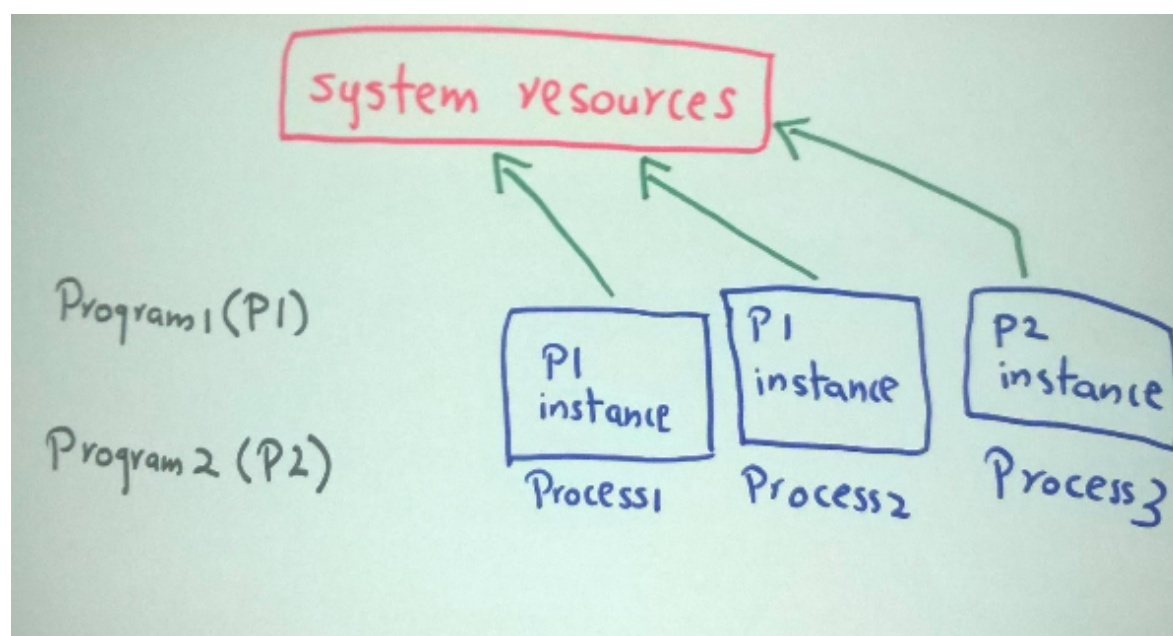
Process Management Activities

- Process creation and termination
- Processes suspension and resumption
- Process synchronization primitives
- Process communication primitives
- Deadlock handling



From Process to Thread

- Single-threaded process has one program counter
 - **program counter** specifies *location of next instruction to execute*
 - processor executes instructions sequentially, one at a time, until completion
- Multi-threaded process has **one program counter per thread**
- **Process is the unit of resource allocation and protection, not thread!**





Resource Management: Memory Management

- Memory is the main storage directly accessible to CPU
 - data needs to be kept in memory before and after processing
 - all instructions should be in memory in order to execute
- Memory management determines what is in memory to **optimize CPU utilization** and **response time, provides a virtual view of memory for programmer**
- Memory management activities:
 - keeping track of which parts of memory are being used and by whom
 - deciding which processes and data to move into and out of memory
 - allocating and deallocating memory space as needed



Resource Management: File Systems

- OS provides a uniform, logical view of data storage
 - **file** is a logical storage unit that abstracts physical properties
 - files are usually organized into **directories**
 - **access control** determines who can access the file
- File system management activities:
 - creating and deleting files and directories
 - primitives to manipulate files and directories
 - mapping files onto secondary storage
 - backup files onto stable (non-volatile) storage media

A special FS: /proc file system

02: Operating System Services & Structures



User Interface

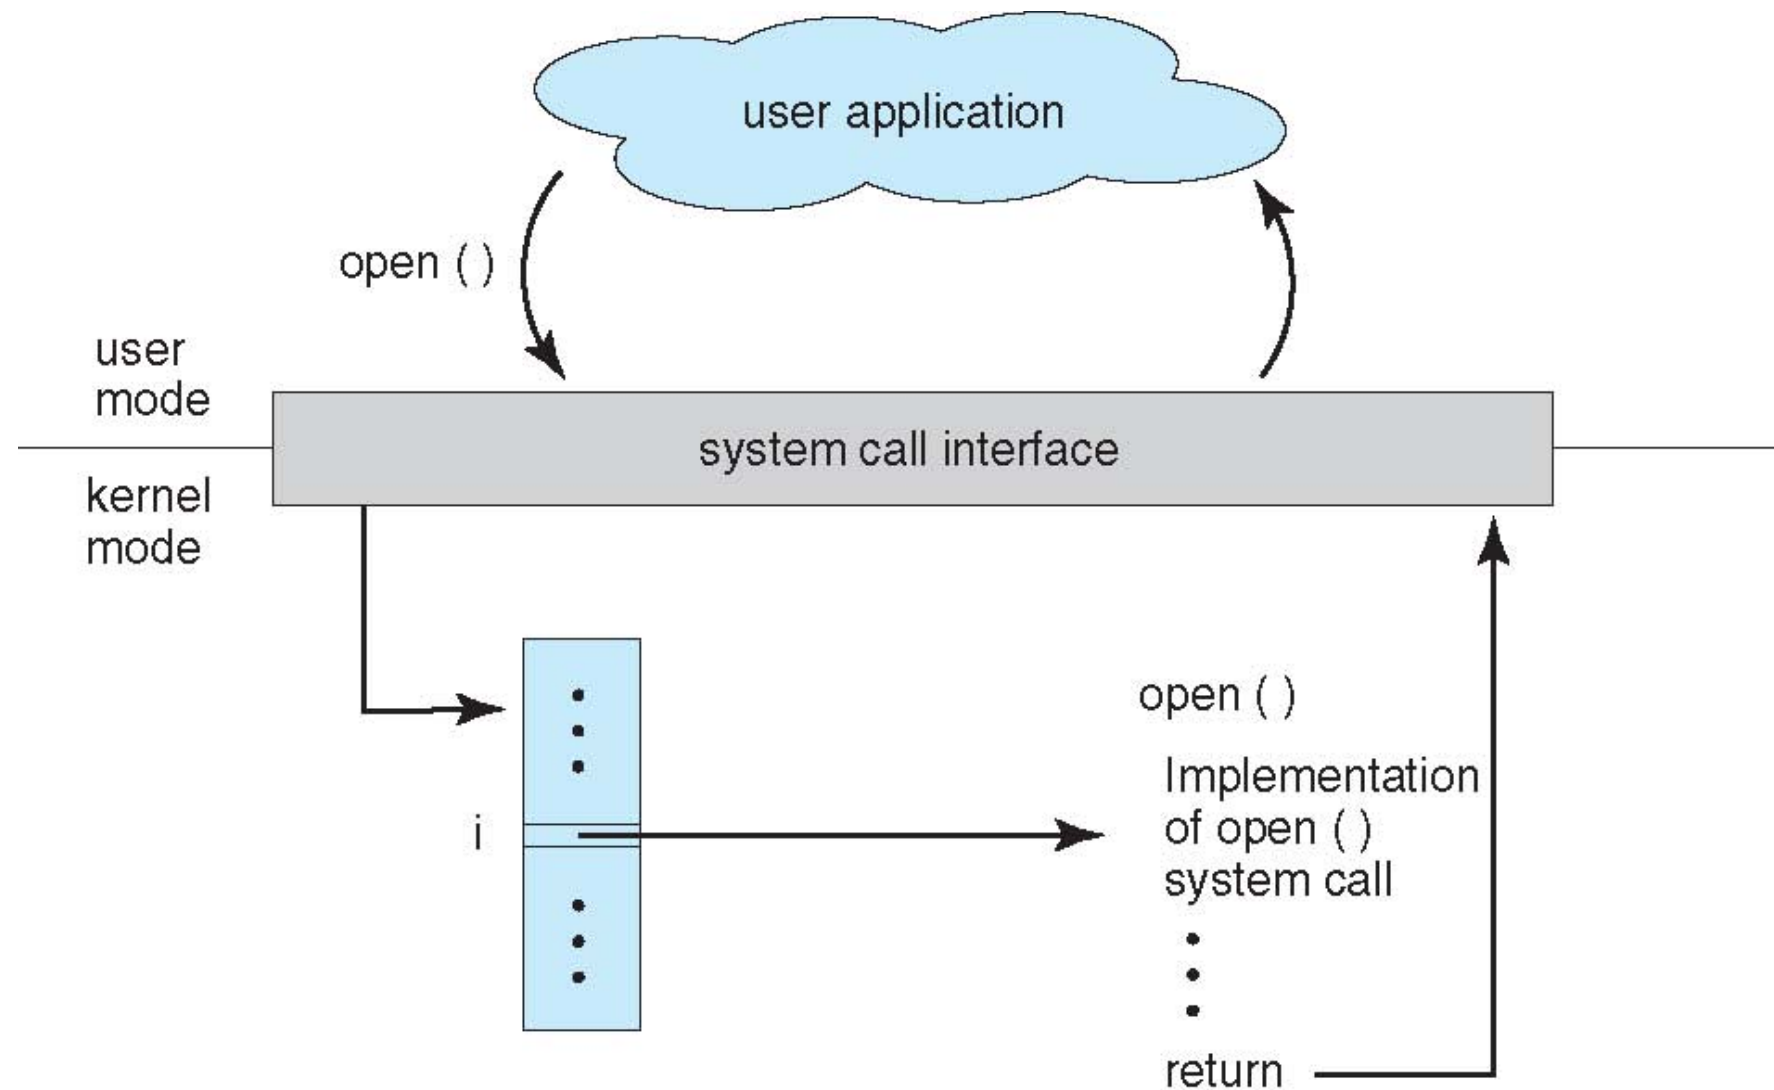
- GUI
- Shell
 - How to chain multiple commands: `cmd1;cmd2;cmd3`
 - `cat`
 - `Find . -name test -print`
 - `Mv`
 - Pipe: `cat xxx | grep xxx`



System Calls

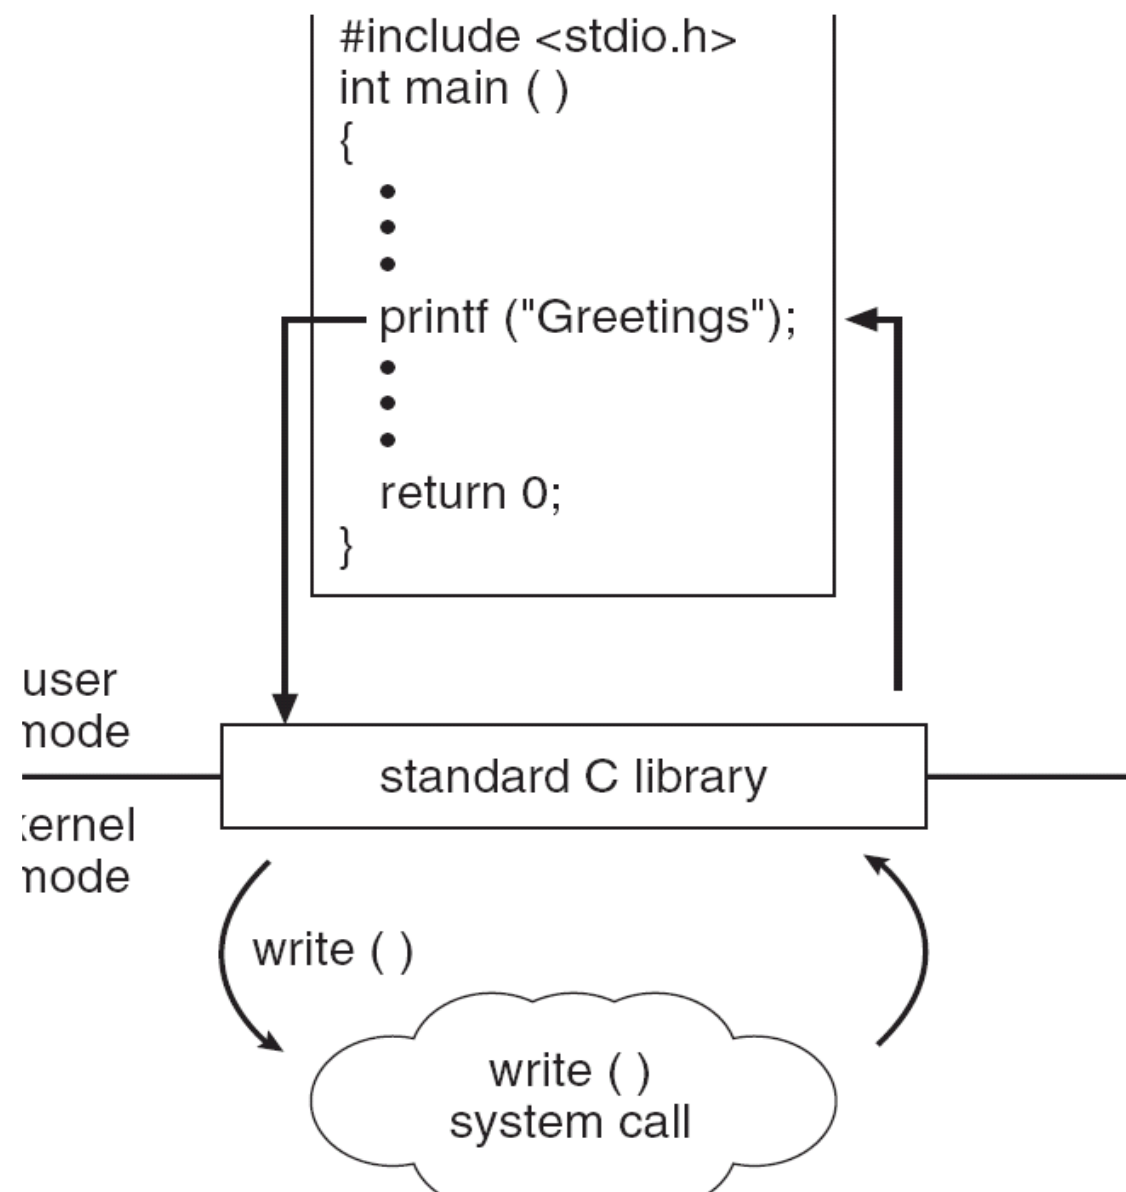
- System call is **a programming interface to access the OS services**
 - Typically written in a high-level language (C or C++)
 - Certain low level tasks are in assembly languages

API – System Call – OS Relationship



Standard C Library Example

- C program invoking printf() library call, which calls write() system call





System Call Parameter Passing

- Parameters are required besides the **system call number**
 - exact type and amount of information vary according to OS and call
- Three general methods to pass parameters to the OS
 - **Register:**
 - pass the parameters in registers
 - simple, but there may be more parameters than registers
 - **Block:**
 - parameters stored in a memory block (or table)
 - address of the block passed as a parameter in a register
 - taken by Linux and Solaris
 - **Stack:**
 - parameters placed, or pushed, onto the stack by the program
 - popped off the stack by the operating system
- Block and stack methods don't limit number of parameters being passed



Operating System Design and Implementation

- Much variation
 - Early OSes in **assembly language**
 - Then system programming languages like Algol, PL/1
 - Now C, C++
- Actually usually a mix of languages
 - Lowest levels in **assembly**
 - Main body in C
 - Systems programs in C, C++, scripting languages like PERL, Python, shell scripts
- More high-level language easier to port to other hardware
 - But slower

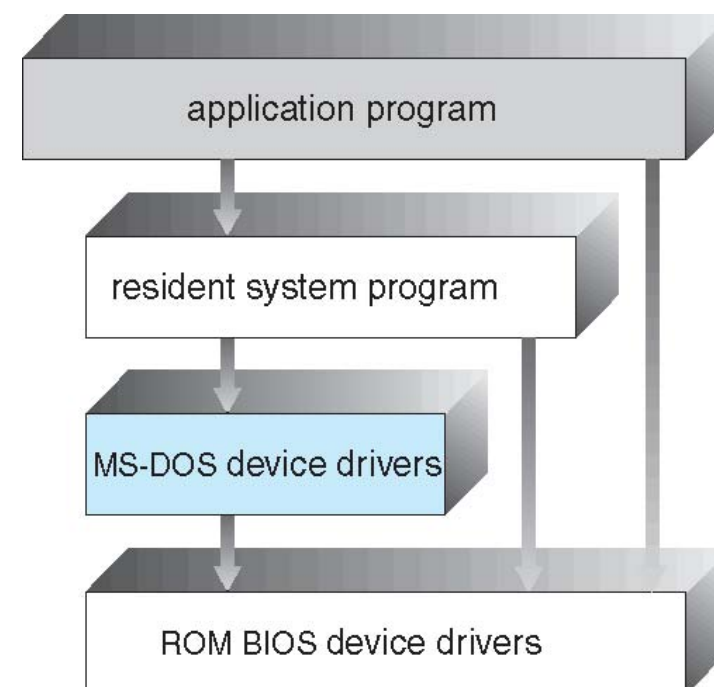


Operating System Structure

- Many structures:
 - **simple structure - MS-DOS**
 - **more complex -- UNIX**
 - **layered structure - an abstraction**
 - **microkernel system structure - L4**
 - **hybrid: Mach, Minix**
 - research system: **exokernel**

Simple Structure: MS-DOS

- No structure at all!: (1981~1994)
 - written to provide the most functionality in the least space
- A typical example: MS-DOS
 - Has some structures:
 - its interfaces and levels of functionality are not well separated
 - the kernel is not divided into modules





Monolithic Structure – Original UNIX

- Limited by hardware functionality, the original UNIX had limited structure
- UNIX OS consists of two separable layers
 - systems programs
 - the kernel: everything below the system-call interface and above physical hardware
 - a large number of functions for one level: file systems, CPU scheduling, memory management ...



Microkernel System Structure

- Microkernel moves as much from the kernel (e.g., file systems) into “user” space
- Communication between user modules uses **message passing**
- Benefits:
 - easier to extend a microkernel
 - **easier to port the operating system to new architectures: code base is small**
 - **more reliable: since less code is running in kernel mode**
 - **more secure: most code is running in unprivileged mode, small code base**
- Detriments:
 - performance overhead of user space to kernel space communication
- Examples: Minix, Mach, QNX, L4...

03: Process



Process Concept

- Process is **a program in execution**, its execution must progress in sequential fashion
 - a program is static and passive, process is dynamic and active
 - **one program can be several processes** (e.g., **multiple instances** of browser, or **even on instance** of the program)
 - process can be started via GUI or command line entry of its name
 - through system calls
- **Process is the basic unit for resource allocation and protection**



Process Concept

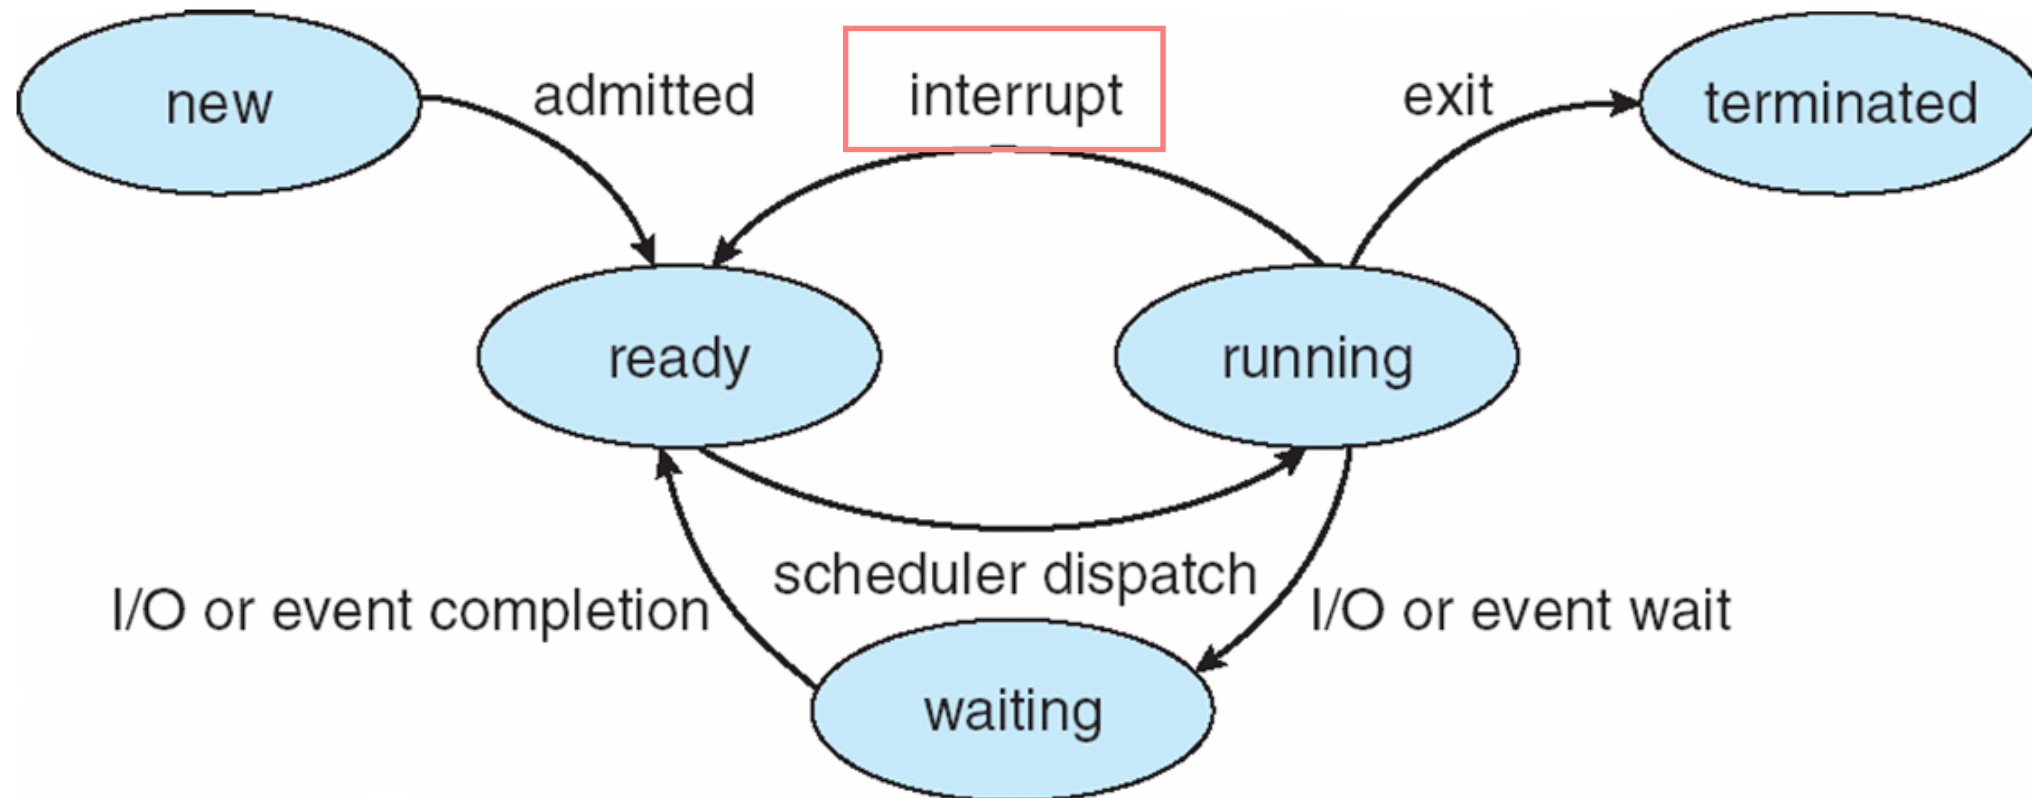
- A process has multiple parts:
 - the program **code**, also called **text section**
 - runtime **CPU states**, including program counter, registers, etc
 - various types of memory:
 - **stack**: temporary data
 - e.g., function parameters, local variables, and *return addresses*
 - **data** section: global variables
 - **heap**: memory dynamically allocated during runtime



Process State

- As a process executes, it changes state
 - **new**: the process is being created
 - **running**: instructions are being executed
 - **waiting/blocking**: the process is waiting for some event to occur
 - **ready**: the process is waiting to be assigned to a processor
 - **terminated**: the process has finished execution

Diagram of Process State





Process Control Block (PCB)

- In the kernel, each process is associated with a **process control block**

- process number (pid)
- process state**
- program counter (PC)**
- CPU registers
- CPU scheduling information
- memory-management data
- accounting data
- I/O status

```
struct task_struct {  
#ifdef CONFIG_THREAD_INFO_IN_TASK  
    /*  
     * For reasons of header soup (see current_thread_info()), this  
     * must be the first element of task_struct.  
     */  
    struct thread_info thread_info;  
#endif  
    /* -1 unrunnable, 0 runnable, >0 stopped: */  
    volatile long state;  
    /*  
     * This begins the randomizable portion of task_struct. Only  
     * scheduling-critical items should be added above here.  
     */  
    randomized_struct_fields_start  
  
    void *stack;  
    atomic_t usage;  
    /* Per task flags (PF_*), defined further below: */  
    unsigned int flags;  
    unsigned int ptrace;
```

- Linux's PCB is defined in struct task_struct: <http://lxr.linux.no/linux+v3.2.35/include/linux/sched.h#L1221>



Threads

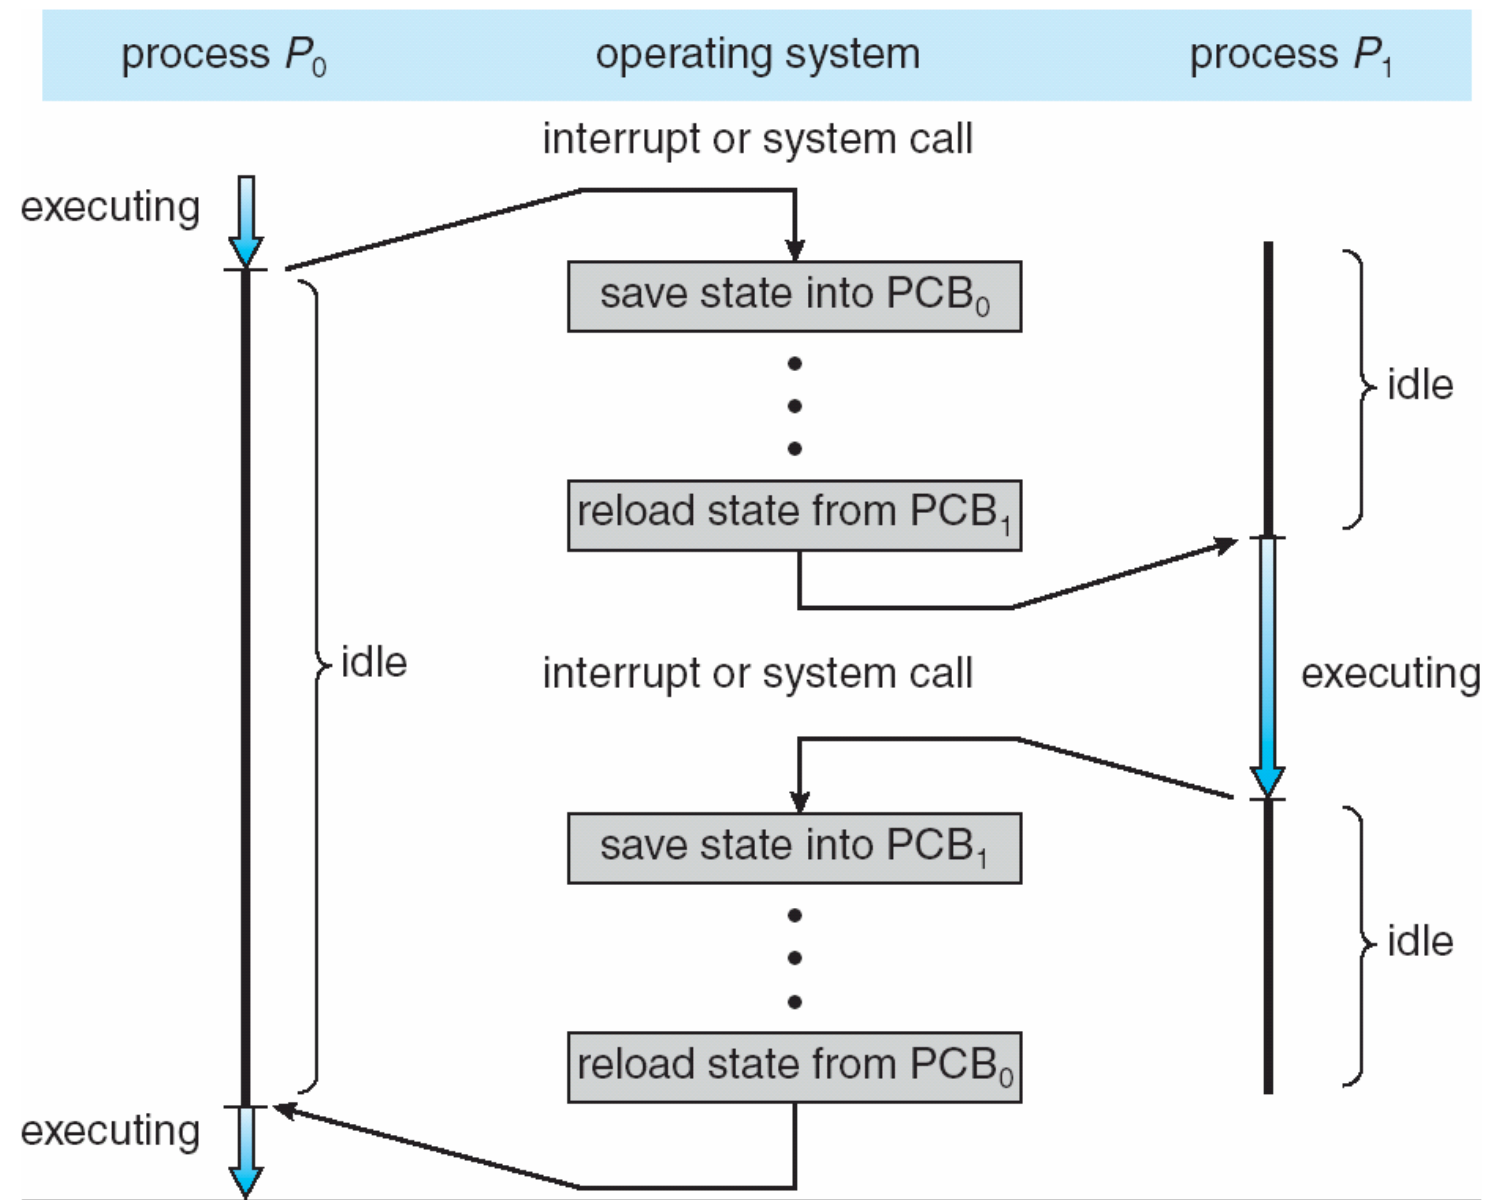
- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - **Multiple locations** can execute at once
 - Multiple threads of control -> threads
- Must then have storage for thread details, multiple program counters in PCB
- Which resources are shared between threads?
 - Stack (no), heap (Y), global data (Y), code (Y)



Context Switch

- **Context switch:** the kernel switches to **another process for execution**
 - save the state of the old process
 - load the saved state for the new process
- **Context-switch is overhead;** CPU does no useful work while switching
 - the more complex the OS and the PCB, longer the context switch
- Context-switch time depends on hardware support
 - some hardware provides multiple sets of registers per CPU: multiple contexts loaded at once

Context Switch





Process Creation

- Parent process creates children processes, which, in turn create other processes, forming **a tree of processes**
 - process identified and managed via a process identifier (pid)
- Design choices:
 - three possible levels of **resource sharing**: all, subset, none
 - parent and children's **address spaces**
 - child duplicates parent address space (e.g., Linux)
 - child has a new program loaded into it (e.g., Windows)
 - **execution** of parent and children
 - **parent and children execute concurrently**
 - parent waits until children terminate



Process Creation

- UNIX/Linux system calls for process creation
 - **fork** creates a new process
 - **exec** overwrites the process' address space with a new program
 - **wait** waits for the child(ren) to terminate



C Program Forking Separate Process

```
#include <sys/types.h>
#include <studio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    pid = fork();                                /* fork another process */
    if (pid < 0) {                                /* error occurred while forking */
        fprintf(stderr, "Fork Failed");
        return -1;
    } else if (pid == 0) {                        /* child process */
        execlp("/bin/ls", "ls", NULL);
    } else {                                      /* parent process */
        wait (NULL);
        printf ("Child Complete");
    }
    return 0;
}
```



fork() multiple times

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    pid = fork();
    If (pid == 0) {
        fork();
    }
    fork();
    return 0;
}
```

How many unique processes are created? (including itself)




Interprocess Communication

- Processes within a system may be independent or cooperating
 - **independent process**: process that cannot affect or be affected by the execution of another process
 - **cooperating process**: processes that can affect or be affected by other processes, including sharing data
 - reasons for cooperating processes: information sharing, computation speedup, modularity, convenience, Security
- Cooperating processes need **interprocess communication** (IPC)
- Two models of IPC
 - **Shared memory**
 - Message passing


Communications Models

(a) Shared memory.



(a)

(b) Message passing.



(b)

Very fast!



POSIX Shared Memory

- POSIX Shared Memory
 - Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```
 - Also used to open an existing segment
 - Set the size of the object: **ftruncate**(shm_fd, 4096);
 - Use **mmap**() to memory-map a file pointer to the shared memory object
 - Reading and writing to shared memory is done by using the pointer returned by **mmap**().



A Simple Kernel Module

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/printk.h>
#include <linux/sched.h>
#include <linux/sched/signal.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Yajin Zhou");
MODULE_DESCRIPTION("A simple example Linux module.");
MODULE_VERSION("0.01");

static int __init os_lkm_example_init(void) {
    struct task_struct *task;

    for_each_process(task)
        printk(KERN_INFO "%s [%d]\n", task->comm, task->pid);

    return 0;
}

static void __exit os_lkm_example_exit(void) {
    printk(KERN_INFO "Goodbye, World!\n");
}

module_init(os_lkm_example_init);
module_exit(os_lkm_example_exit);
```

```
/*
 * hello-1.c - The simplest kernel module.
 */
#include <linux/module.h>      /* Needed by all modules */
#include <linux/kernel.h>      /* Needed for KERN_INFO */

int init_module(void)
{
    printk(KERN_INFO "Hello world 1.\n");

    /*
     * A non 0 return means init_module failed; module can't be loaded.
     */
    return 0;
}

void cleanup_module(void)
{
    printk(KERN_INFO "Goodbye world 1.\n");
}
```

insmod xxx.ko

Make menuconfig