



# So Far...

- ▶ Our goal (supervised learning)
- ▶ Learn a set of discriminant functions
  - Bayesian framework
    - We could design an optimal classifier if we knew:
      - $P(\omega_i)$  : priors and  $P(x | \omega_i)$  : class-conditional densities
      - Using training data to estimate  $P(\omega_i)$  and  $P(x | \omega_i)$
  - Directly learning discriminant functions from the training data (assume the form of the function is known)
    - Linear Regression
    - Logistic Regression
    - SVM
    - Kernel methods
    - Perceptron
    - Neural Network
- ▶ Other possible approaches?

# k Nearest Neighbor Classifier

**Deng Cai (蔡登)**

College of Computer Science  
Zhejiang University

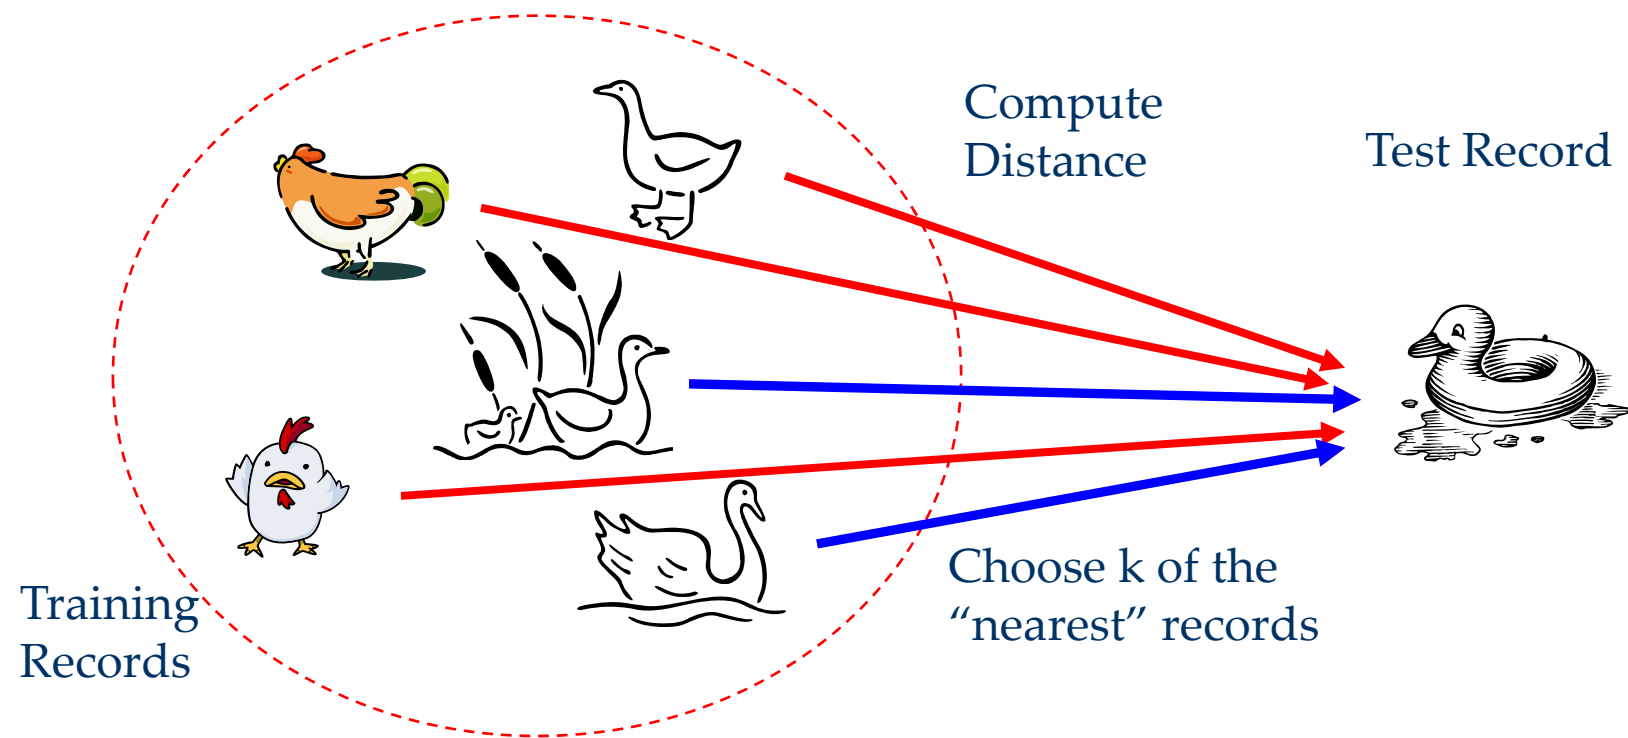
dengcai@gmail.com






# Nearest Neighbor Classifiers

- ▶ Basic idea:
  - If it walks like a duck, quacks like a duck, then it's probably a duck






# Nearest-Neighbor Classifiers




- Requires three things
  - The set of stored records
  - Distance Metric to compute distance between records
  - The value of  $k$ , the number of nearest neighbors to retrieve
- To classify an unknown record:
  - Compute distance to other training records
  - Identify  $k$  nearest neighbors
  - Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)




# Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

K-nearest neighbors of a record  $x$  are data points that have the  $k$  smallest distance to  $x$



# How many parameters in kNN?

- ▶ A Linear Classifier

$$f(x) = w^T x$$


- The number of parameters?

- ▶ kNN Classifier

- The number of parameters?




# 15-Nearest Neighbor Classifier





# 1-Nearest Neighbor Classifier







# How many parameters in kNN?

- ▶ A Linear Classifier

$$f(x) = w^T x$$

- The number of parameters?


- ▶ kNN Classifier

- **Effective** number of parameters?

$$\frac{N}{k}$$



# 1 nearest-neighbor



Voronoi diagram  
(tessellation)

# Decision Tree

**Deng Cai (蔡登)**

College of Computer Science  
Zhejiang University

dengcai@gmail.com






# Decision Tree

## ► Yes/No Question Game

- Aim: To guess the name of a famous person or character by asking yes/no questions
- Example questions:

- Are you male?
- Are you a real person?
- Are you an actor?
- Are you alive?
- Have you won an Oscar?
- Do you play a sport?
- Are you an athlete?
- Do you come from the USA?
- Do you work in the entertainment industry?
- Do you play a musical instrument?
- Are you Japanese?
- Are you a scientist?
- Are you a cartoon character?






# Decision Tree

- ▶ A hierarchical data structure that represents data by implementing a divide and conquer strategy
  - Given a collection of examples, learn a decision tree that represents it.
  - Use this representation to classify new examples
- ▶ Can be used as a non-parametric classification and regression method.



# Decision Tree





# Decision Tree

- ▶ Decision Trees are classifiers for instances represented as features vectors (color= ;shape= ;label= )
- ▶ Nodes are tests for feature values;
- ▶ There is one branch for each value of the feature
- ▶ Leaves specify the categories (labels)
- ▶ Can categorize instances into multiple disjoint categories
- ▶ Output is a discrete category. Real valued outputs are possible (regression trees)
- There are efficient algorithms for processing large amounts of data. (But not too many features)
- There are methods for handling noisy data (classification noise and attribute noise) and for handling missing attribute values.



# Basic Decision Tree Learning Algorithm

- ▶ Data is processed in Batch (I.e., all the data is available).
- ▶ Recursively build a decision tree top-down.
- DT(Examples, Attributes)
  - If all Examples have same label:  
return a leaf node with Label
  - Else
    - If Attributes is empty:  
return a leaf with majority Label
    - Else
      - Pick an attribute A as root
      - For each value v of A
      - Let Examples(v) be all the examples for which A=v
      - Add a branch out of the root for the test A=v
        - If Examples(v) is empty  
create a leaf node labeled with the majority label in Examples
        - Else  
recursively create subtree by calling DT(Examples(v), Attribute-{A})





# Decision Tree

## ► Yes/No Question Game

- Aim: To guess the name of a famous person or character by asking yes/no questions
- Example questions:
  - Are you male?
  - Are you a real person?
  - Are you an adult?
  - Are you alive?
  - Have you written a famous book?
  - Do you play a sport?
  - Are you an actor?
  - Do you come from England?
  - Do you work in Hollywood?
  - Do you play a musical instrument?
  - Are you Japanese?
  - Are you a scientist?
  - Are you a cartoon character?

How many trees we can build?

All the trees are the same? Or  
Some trees are better than  
others?



# Decision Tree

- ▶ A better tree (model)

- Small trees

$$\text{EPE}(f) = (\text{bias})^2 + \text{variance} + \text{noise}$$

- Same bias, less variance.



- ▶ Finding the minimal decision tree **consistent with the data** is NP-hard
- ▶ The recursive algorithm is a greedy heuristic search for a simple tree, but cannot guarantee optimality.
- ▶ The main decision in the algorithm is the selection of the next attribute to condition on.



# Picking the Best Feature (Attribute)

- Consider data with two Boolean attributes (A,B).
  - $\langle (A=0, B=0), - \rangle$ : 50 examples
  - $\langle (A=0, B=1), - \rangle$ : 50 examples
  - $\langle (A=1, B=0), - \rangle$ : 0 examples
  - $\langle (A=1, B=1), + \rangle$ : 100 examples

- What should be the first attribute we select? A or B?
- Splitting on A**: we get purely labeled nodes
- Splitting on B**: we don't get purely labeled nodes.






# Picking the Best Feature (Attribute)

- Consider data with two Boolean attributes (A,B).
  - $\langle (A=0, B=0), - \rangle$ : 50 examples
  - $\langle (A=0, B=1), - \rangle$ : 50 examples
  - $\langle (A=1, B=0), - \rangle$ : **3** examples
  - $\langle (A=1, B=1), + \rangle$ : 100 examples

- Trees looks structurally similar; which attribute should we choose?



Advantage A. But...  
Need a way to quantify things



# Picking the Best Feature (Attribute)

- ▶ We want attributes that split the examples to sets that are relatively pure in one label; this way we are closer to a leaf node.
- ▶ The most popular heuristics is based on information gain, originated with the ID3 system of Quinlan.



# Entropy

- ▶ Entropy (impurity, disorder) of a set of examples,  $S$ , relative to a binary classification is:

$$Entropy(S) = -P_+ \log P_+ - P_- \log P_-$$


- where  $P_+$  is the proportion of positive examples in  $S$
- $P_-$  is the proportion of negative examples
- If all the examples belong to the same category  $Entropy = 0$
- If the examples are equally mixed (0.5,0.5)  $Entropy = 1$
- ▶ In general, when  $p_i$  is the fraction of examples labeled  $i$ :

$$Entropy(S) = - \sum_{i=1}^c P_i \log P_i$$

- ▶ Entropy can be viewed as the number of bits required, on average, to encode the class of labels. If the probability for + is 0.5, a single bit is required for each example; if it is 0.8 -- can use less than 1 bit.



# Entropy





# Picking the Best Feature (Attribute)

- ▶ We want attributes that split the examples to sets **that are relatively pure in one label**; this way we are closer to a leaf node.
- ▶ We can pick the feature that the resulting data partitions have low entropy





# Information Gain

- ▶ **The information gain** of an attribute  $a$  is the expected reduction in entropy caused by partitioning on this attribute.

$$Gain(S, a) = Entropy(S) - \sum_{v \in values(a)} \frac{|S_v|}{|S|} Entropy(S_v)$$


- where  $S_v$  is the subset of  $S$  for which attribute  $a$  has value  $v$
- ▶ Partitions of low entropy lead to high gain.



# Picking the Best Feature (Attribute)

- Consider data with two Boolean attributes (A,B).
  - $\langle (A=0, B=0), - \rangle$ : 50 examples
  - $\langle (A=0, B=1), - \rangle$ : 50 examples
  - $\langle (A=1, B=0), - \rangle$ : **3** examples
  - $\langle (A=1, B=1), + \rangle$ : 100 examples

- Trees looks structurally similar; which attribute should we choose?



The information gain of A and B



# An Illustrative Example

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Weak	Yes
4	Rain	Mild	High	Weak	Yes
5	Rain	Cool	Normal	Weak	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
10	Rain	Mild	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Weak	Yes
14	Rain	Mild	High	Strong	No



## An Illustrative Example (2)

$$\begin{aligned} & \text{Entropy}(S) \\ &= -\frac{9}{14} \log\left(\frac{9}{14}\right) \\ &\quad - \frac{5}{14} \log\left(\frac{5}{14}\right) \\ &= 0.94 \end{aligned}$$

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Weak	Yes
4	Rain	Mild	High	Weak	Yes
5	Rain	Cool	Normal	Weak	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
10	Rain	Mild	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Weak	Yes
14	Rain	Mild	High	Strong	No

9+,5-



# An Illustrative Example (2)

		Humidity	Wind	PlayTennis	
		High	Weak	No	
		High	Strong	No	
		High	Weak	Yes	
		High	Weak	Yes	
		Normal	Weak	Yes	
		Normal	Strong	No	9+,5-
		Normal	Strong	Yes	E=.94
		High	Weak	No	
		Normal	Weak	Yes	
		Normal	Weak	Yes	
		Normal	Strong	Yes	
		High	Strong	Yes	
		Normal	Weak	Yes	
		High	Strong	No	

Humidity	Wind
High	Weak
Normal	Strong
Weak	Weak
Strong	Strong

3+,4- E=.985	6+,1- E=.592	6+2- E=.811	3+,3- E=1.0
-----------------	-----------------	----------------	----------------

Gain(S, Humidity) = .94 - 7/14 * 0.985 - 7/14 * 0.592 = 0.151	Gain(S, Wind) = .94 - 8/14 * 0.811 - 6/14 * 1.0 = 0.048
---	---

$$Gain(S, a) = Entropy(S) - \sum_{v \in values(a)} \frac{|S_v|}{|S|} Entropy(S_v)$$



## An Illustrative Example (3)

$$Gain(S, Humidity) = 0.151$$

$$Gain(S, Wind) = 0.048$$

$$Gain(S, Temperature) = 0.029$$


$$Gain(S, Outlook) = \mathbf{0.246}$$

Outlook

A diagram of a decision tree node. The word "Outlook" is positioned above a vertex. Three lines radiate from this vertex: one pointing up and to the left, one pointing straight down, and one pointing up and to the right.




## An Illustrative Example (3)



Day	Outlook	PlayTennis
1	Sunny	No
2	Sunny	No
3	Overcast	Yes
4	Rain	Yes
5	Rain	Yes
6	Rain	No
7	Overcast	Yes
8	Sunny	No
9	Sunny	Yes
10	Rain	Yes
11	Sunny	Yes
12	Overcast	Yes
13	Overcast	Yes
14	Rain	No



## An Illustrative Example (3)



Continue until:

- Every attribute is included in **path**, or,
- All examples in the leaf have same label

Day	Outlook	PlayTennis
1	Sunny	No
2	Sunny	No
3	Overcast	Yes
4	Rain	Yes
5	Rain	Yes
6	Rain	No
7	Overcast	Yes
8	Sunny	No
9	Sunny	Yes
10	Rain	Yes
11	Sunny	Yes
12	Overcast	Yes
13	Overcast	Yes
14	Rain	No



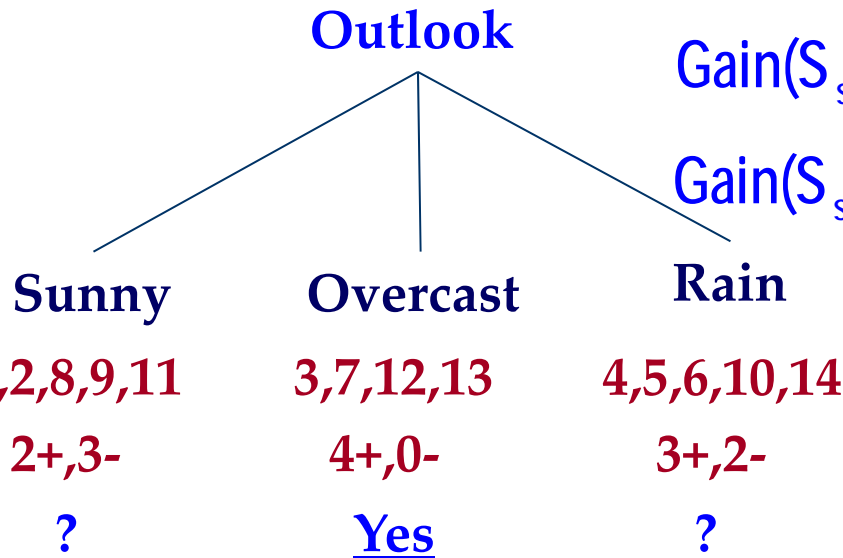


# An Illustrative Example (4)

$$\text{Gain}(S_{\text{sunny}}, \text{Humidity}) = .97 - (3/5) 0 - (2/5) 0 = .97$$

$$\text{Gain}(S_{\text{sunny}}, \text{Temp}) = .97 - 0 - (2/5) 1 = .57$$

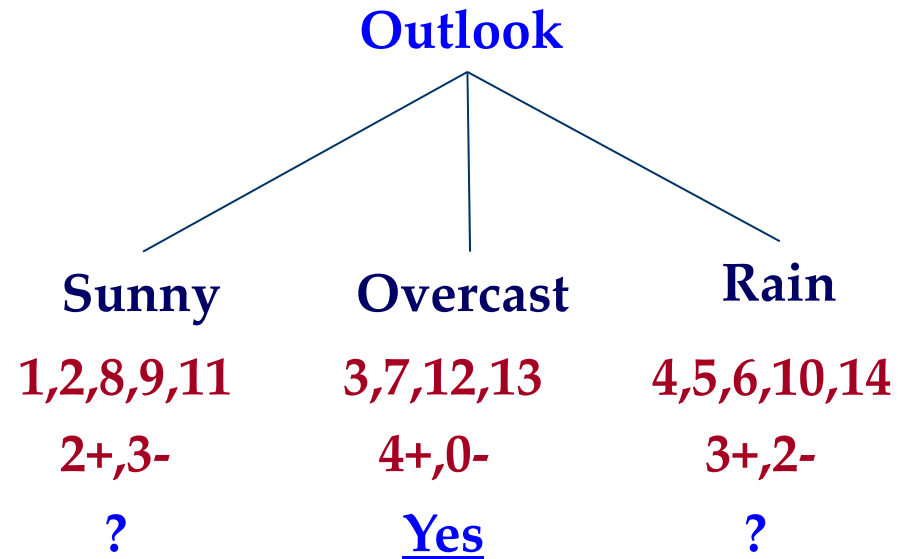
$$\text{Gain}(S_{\text{sunny}}, \text{Wind}) = .97 - (2/5) 1 - (3/5) .92 = .02$$



Day	Outlook	Temperature	Humidity	Wind	PlayTennis
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes

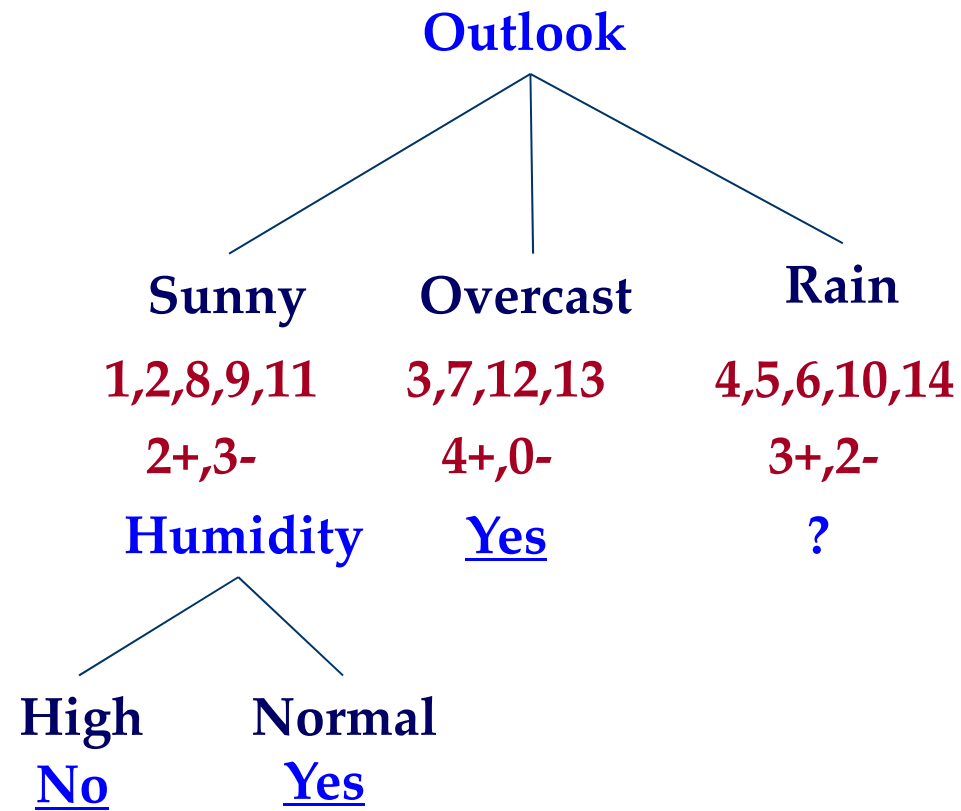


## An Illustrative Example (5)






## An Illustrative Example (5)





## An Illustrative Example (6)






# Summary: ID3(Examples, Attributes, Label)

- Let  $S$  be the set of Examples  
     $Label$  is the target attribute (the prediction)  
     $Attributes$  is the set of measured attributes
  - 
  - Create a Root node for tree
  - If all examples are labeled the same return a single node tree with  $Label$
  - Otherwise Begin
    - $A$  = attribute in  $Attributes$  that best classifies  $S$
    - for each possible value  $v$  of  $A$ 
      - Add a new tree branch corresponding to  $A=v$
      - Let  $S_v$  be the subset of examples in  $S$  with  $A=v$
      - if  $S_v$  is empty: add leaf node with the common value of  $Label$  in  $S$
      - Else: below this branch add the subtree
        - $ID3(S_v, Attributes - \{a\}, Label)$
  - End
- Return Root



# History of Decision Tree Research

- ▶ Hunt and colleagues in Psychology used full search decision trees methods to model human concept learning in the 60's
- ▶ Quinlan developed ID3, with the information gain heuristics in the late 70's to learn expert systems from examples
- ▶ Breiman, Friedmans and colleagues in statistics developed CART (classification and regression trees) simultaneously
- ▶ A variety of improvements in the 80's: coping with noise, continuous attributes, missing data, non-axis parallel etc.
- ▶ Quinlan's updated algorithm, C4.5 (1993) is commonly used (New:C5)
- ▶ Boosting (or Bagging) over DTs is a very good general purpose algorithm

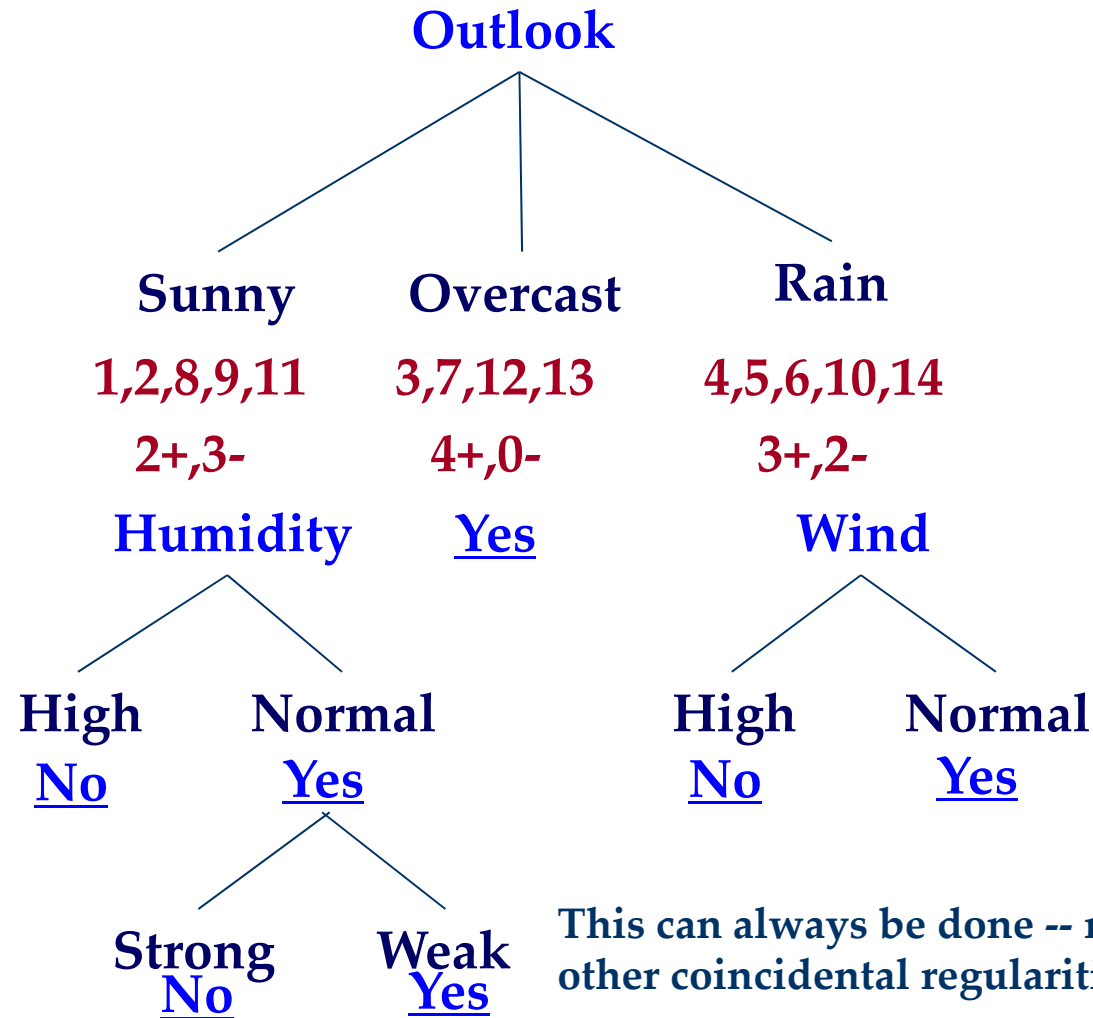


$$\text{EPE}(f) = (\text{bias})^2 + \text{variance} + \text{noise}$$



# Decision Trees: Low Bias Model

Outlook = Sunny, Temp = Hot, Humidity = Normal, Wind = Strong, **NO**



This can always be done -- may fit noise or other coincidental regularities





# Avoid Overfitting

- ▶ How the overfitting can this be avoided with linear classifiers?
- ▶ For Decision Trees: Two basic approaches
  - Prepruning:
    - Stop growing the tree at some point during construction when it is determined that there is not enough data to make reliable choices.
  - Postpruning:
    - Grow the full tree and then remove nodes that seem not to have sufficient evidence.
- ▶ Validation Set